

Anudeep Manda

mandal@purdue.edu | (925) 683-7968 | [linkedin.com/in/anudeepmanda](https://www.linkedin.com/in/anudeepmanda)

Dear OpenAI Applied Engineering Team,

I am writing to express my strong interest in the Software Engineering Intern (Emerging Talent) role on the Applied Engineering team at OpenAI. I am a Computer Engineering and Applied Mathematics student at Purdue University, and I am deeply motivated by OpenAI's mission to responsibly deploy AI systems that are useful, reliable, and accessible to millions of users worldwide. OpenAI's emphasis on learning from real-world deployment and building high-quality, user-centered AI products strongly aligns with how I approach engineering.

In my recent role as a Software Engineering Intern at USG Manufacturing, I owned the development of high-throughput backend services and real-time analytics pipelines processing over 60 RTSP camera feeds under strict latency and reliability constraints. I built backend APIs in Python and SQL, developed React-based dashboards for operational visibility, and deployed containerized GPU-backed inference services using Docker and CI/CD pipelines. This work required careful reasoning about system performance, correctness, and production reliability—skills directly applicable to building and scaling customer-facing AI products at OpenAI.

Previously, at an early-stage AI startup, I designed and productionized backend services using Python and TypeScript to support AI-powered, customer-facing features. I worked across the full stack, integrating React-based UIs with backend services and databases, transforming unstructured data into structured, queryable formats for downstream analytics and AI workflows. I regularly profiled latency, handled failure modes, and iterated quickly based on user feedback—experience that mirrors the fast-paced, iterative environment of OpenAI's Applied teams.

Beyond my internships, my projects reflect a strong interest in applied AI systems and product engineering. These include a Cloudflare-inspired bot detection proxy focused on high-throughput request ingestion, observability, and secure APIs, as well as a domain-specific language model for interpreting California Tax Code queries, where correctness, edge-case handling, and safe user-facing behavior were critical. Across my work, I prioritize building systems that are robust, measurable, and designed to improve with real-world usage.

I am particularly excited about the opportunity to contribute to OpenAI's efforts to bring AI safely and responsibly into real-world products like ChatGPT and the OpenAI API. I would be thrilled to work alongside engineers, researchers, and product teams to build scalable, user-focused AI experiences that have a meaningful global impact. Thank you for your time and consideration—I would welcome the opportunity to further discuss how my background and interests align with OpenAI's Applied Engineering team.

Sincerely,
Anudeep Manda